

FEATURE LIST

1. Voice Assistant (LiveKit Agent)

- ✓ Real-time **voice interaction** with users via LiveKit.
 - ✓ Integrated **Google Gemini LLM** for conversational responses.
 - ✓ **Text-to-Speech (TTS)** powered by Gemini with configurable voice.
 - ✓ Transcription filtering to ignore short or filler words (e.g., *uh*, *um*, *okay*).
 - ✓ Friendly, conversational, human-like response style (designed for store staff feel).
 - ✓ Integrated **Retrieval-Augmented Generation (RAG)** for store knowledge lookup.
 - ✓ Avoids repeating the same knowledge chunks in one conversation.
 - ✓ Latency logging for both search queries and overall response time.
 - ✓ Greets the user automatically at the start of the conversation.
 - ✓ Function tool (`info_search`) exposed to the LLM for knowledge searches.
 - ✓ Seamless backend logic for handling room connections, session lifecycle, and input options.
-

2. RAG Knowledge Base Indexing System

- ✓ Loads all PDFs from a specified `data/` folder for knowledge ingestion.
 - ✓ Splits large documents into manageable chunks for better search and recall.
 - ✓ Configurable chunk size and overlap to preserve context across splits.
 - ✓ Automatically generates **unique IDs** for each knowledge chunk using UUIDs.
 - ✓ Stores all document chunks as embeddings in a Qdrant vector database.
 - ✓ Uses **OpenAI's text-embedding-3-small** model for generating high-quality embeddings.
 - ✓ Optional `--reset` flag to wipe and recreate the Qdrant collection.
 - ✓ Qdrant collection created with 1536-dimension vectors and cosine distance metric for similarity search.
 - ✓ Modular, clean design makes it easy to swap embedding models (Ollama, Bedrock, OpenAI).
 - ✓ Logs the number of chunks being added to Qdrant for visibility.
-

3. Embedding Model Flexibility

- ✓ Central `get_embedding_function` for managing embedding provider selection.
- ✓ Supports OpenAI embeddings out of the box.

- ✓ Provides easy-to-activate alternatives for Ollama (local) and AWS Bedrock embeddings.
 - ✓ Environment variable-based API key handling for secure setup.
-

4. Additional System Features

- ✓ Environment variable management with `.env` for secure configuration.
- ✓ Clear, consistent logging throughout indexing and agent processes.
- ✓ Clean separation of concerns: indexing, embedding logic, and agent behavior are modular.
- ✓ Prepared for expansion with other LiveKit plugins (silero, cartesia, etc., based on your previous setup).